




GazeEarth

Input evidence request

UHR RS Image ($10K \times 10K$ px)



Q: How many mines are there in the entire map?

- A. 8 B. 6
C. 1. D. 7

Resize


1536×1536 px




Stage 1: Question guided gaze



Question determines what deserve pixels




4x4 indexed overview
R remains bound to I



Read question and overview,
identify sufficient evidence
without answering yet

Selected Region on the Overview




Region remains bound to the full image

status: selected
region_ids: [I1:c05]
spatial pointer Region

Return a
spatial pointer


Stage 2: Deterministic topology-preserving focus

1. Bind the gaze to the scene



Reallocate pixels within one canvas—not
by opening another crop.


2. Reallocate a fixed pixel budget



Reallocate pixels within one canvas, not by opening another crop.




3. Render a single visual field I^*



Stage 3: Evidence-grounded answer

Answer from one question-conditioned visual field.

Resized Image $I^* \times 2048 \times 2048$ px



Same frozen MLLM

Answer

Convert the
question into a
compact spatial
request.

Alright, I first inspect the fil $10K \times 10k$ scene to locate where the counting evidence is concentrated. ... The overview points to one candidate open pit mine in region C05, while the remaining regions mainly contain farmland, settlements, roads, and rivers. ... I therefore gaze at C05 and allocate more pixels to its stepped excavation and surrounding pit, while retaining the whole map for spatial context. ... The focused view confirms one distinct mine. final Answer: 1 ✓